

Chi Zhang

✉ iskyzh@gmail.com | 🌐 skyzh.dev | 🐙 [skyzh](https://github.com/skyzh) | 🌐 [alex-chi-skyzh](https://www.linkedin.com/in/alex-chi-skyzh)

Education

Carnegie Mellon University

Master of Science in Computer Science

August 2022(Expected) – December 2023 (Expected)

Pittsburgh, PA

Shanghai Jiao Tong University

B.Eng in Computer Science and Technology

September 2018 – June 2022

Shanghai, China

- GPA 93.58/100, Rank 1/149, National Scholarship 2019 (Top 0.2% national-wide)
 - A+ Courses: Operating Systems, Computer Architecture, Computer Networks, and 28 others
-

Internship Experience

Singularity Data, Inc.

Database System R&D Intern

August 2021 – July 2022

Shanghai, China

- One of the top contributors of 🐙 [RisingWave](https://github.com/risingwave) — a next-generation streaming database in the cloud. Worked on the development of almost all components related to stream computing and state store.
- Designed and implemented **shared state** to support **streaming index** in RisingWave; implemented **lookup join executor** based on shared state to support efficient **index delta joins**.
- Lead the team to investigate and analyze performance issues in RisingWave with benchmarks; fixed bugs and proposed strategies which improved the system throughput by 10x in a 3-month period.
- Greatly improved RisingWave's development experience by initiating the developer ecosystem, including streaming system dashboard, developers' tool RiseDev, and a benchmark set-up tool based on Terraform. They are now indispensable parts of everyone's development process.

ByteDance, Ltd.

Storage System R&D Intern, TerarkDB Team

June 2021 – August 2021

Beijing, China

- Implemented **Zone-Aware Garbage Collection** in **TerrakDB** for Zoned Namespace SSDs, which reduced 3-4x of space amplification caused by interleaving write lifetime in a single ZNS zone. 🐙
- Added observability facilities to **ZenFS** (by Western Digital) to analyze bottlenecks and implemented a **WAL-Aware Zone Allocator**, which reduced the p999 tail latency by 100x. 🐙

PingCAP, Inc.

Sotrage System R&D Intern, TiKV Storage Team

August 2020 – January 2021

Shanghai, China

- Built LSM-based storage engine **AgateDB** from ground-up. Inspired by WiscKey and BadgerDB, AgateDB separates large vallues from LSM tree into value log, so as to reduce write amplification. 🐙
-

Open-Source Contributions

RisingLight Community

RisingLight Project Maintainer

January 2022 – Now

🐙 [risinglightdb](https://github.com/risinglightdb)

- Leads the development of **RisingLight**, an OLAP database system for educational purpose.
- Designed and implemented a merge-tree **column storage engine** with snapshot isolation support; mentored community members to add functionalities like **filter scan** and **statistics**.

TiKV Community

TiKV Maintainer

May 2020 – Now

🐙 [tikv](https://github.com/tikv)

- Maintains **TiKV Coprocessor**, the push-down execution framework of TiDB.
- Mentored community members to contribute new feature to TiKV in the **LFX Mentorship**. 🐙 🐙

Personal Projects

4k followers 🐙 [skyzh](https://github.com/skyzh)

- **type-exercise-in-rust** (☆ 870): Learn Rust black magics (GAT, HRTB, bypassing compiler bugs, macros) by implementing an expression framework in database systems. 🐙
-

Skills

Programming Languages: Rust, C++, Golang, Python and Node.js

Tech Skills: Key-Value Storage Systems, SSD-optimized File Systems, Database Systems, Stream-Processing Systems

Updated on August 25, 2022